

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder

# Reload the dataset
df = pd.read_csv("customer_purchase_prediction.csv")
df.head()

```

	Age	Annual_Income	Gender	Marital_Status	Browsing_Time \
0	56	21920	Male	Married	4.589678
1	69	126121	Female	Divorced	14.026686
2	46	97219	Female	Married	9.317701
3	32	96872	Female	Single	12.214115
4	60	101132	Female	Married	6.560902

	Previous_Purchases	Clicked_Ad	Customer_Rating	Purchase
0	17	0	2.070361	0
1	5	1	1.610790	0
2	8	1	4.073605	1
3	5	0	2.744720	0
4	17	0	1.708211	0

```

# Encode categorical variables
label_encoder = LabelEncoder()
df["Gender"] = label_encoder.fit_transform(df["Gender"]) # Male: 1,
Female: 0
df["Marital_Status"] =
label_encoder.fit_transform(df["Marital_Status"])

# Define features (X) and target (y)
X = df.drop(columns=["Purchase"])
y = df["Purchase"]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Initialize and train the decision tree classifier
classifier = DecisionTreeClassifier(random_state=42)
classifier.fit(X_train, y_train)

# Make predictions
y_pred = classifier.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)

```

```

accuracy, classification_rep
(1.0,
 1.00      precision      recall  f1-score      support\n\n
0      1.00      1.00      1.00      920\n      1      1.00
1.00      1.00      80\n\n      accuracy
1.00      1000\n      macro avg      1.00      1.00      1.00      1000\n
nweighted avg      1.00      1.00      1.00      1000\n')

from sklearn.model_selection import cross_val_score

# Perform cross-validation to evaluate the model's performance on
unseen data
cv_scores = cross_val_score(classifier, X, y, cv=5,
scoring='accuracy')

# Calculate mean and standard deviation of cross-validation scores
cv_mean = np.mean(cv_scores)
cv_std = np.std(cv_scores)

cv_mean, cv_std, cv_scores

(0.9996, 0.0004898979485566361, array([1.    , 1.    , 0.999, 0.999, 1.
]))

```

USING UNSEEN DATA

```

# Generate unseen data with similar characteristics as the original
dataset
n_unseen = 5000 # Number of rows for unseen data

unseen_data = {
    "Age": np.random.randint(18, 70, size=n_unseen),
    "Annual_Income": np.random.randint(20000, 150000, size=n_unseen),
    "Gender": np.random.choice(["Male", "Female"], size=n_unseen),
    "Marital_Status": np.random.choice(["Single", "Married",
"Divorced", "Widowed"], size=n_unseen),
    "Browsing_Time": np.random.uniform(1, 15, size=n_unseen),
    "Previous_Purchases": np.random.randint(0, 20, size=n_unseen),
    "Clicked_Ad": np.random.choice([0, 1], size=n_unseen, p=[0.7,
0.3]),
    "Customer_Rating": np.random.uniform(1, 5, size=n_unseen),
}

# Add the target variable "Purchase" using similar logic
unseen_data["Purchase"] = np.where(
    (unseen_data["Browsing_Time"] > 5) &
    (unseen_data["Annual_Income"] > 50000) &
    (unseen_data["Previous_Purchases"] > 2) &

```

```

    (unseen_data["Customer_Rating"] > 3) &
    (unseen_data["Clicked_Ad"] == 1),
    1,
    0
)

# Convert to a DataFrame
unseen_df = pd.DataFrame(unseen_data)

# Save unseen data to CSV
unseen_df.to_csv('unseen_df.csv', index=False)
unseen_df

```

	Age	Annual_Income	Gender	Marital_Status	Browsing_Time	\
0	19	29964	Female	Single	3.740645	
1	22	130043	Female	Single	14.909457	
2	54	44838	Female	Single	9.805637	
3	41	102474	Female	Married	2.083550	
4	22	30597	Female	Widowed	11.168576	
...	
4995	41	87046	Male	Divorced	14.117459	
4996	33	120775	Male	Married	9.392444	
4997	63	98042	Male	Married	5.628072	
4998	39	87968	Female	Single	9.236977	
4999	48	95977	Female	Married	11.941578	

	Previous_Purchases	Clicked_Ad	Customer_Rating	Purchase
0	0	0	1.429187	0
1	17	0	3.285423	0
2	9	0	4.362337	0
3	17	0	1.598685	0
4	15	0	2.754389	0
...
4995	7	0	3.181555	0
4996	15	1	2.078366	0
4997	7	1	3.615173	1
4998	5	1	4.616142	1
4999	19	0	1.781354	0

```

[5000 rows x 9 columns]

# Prepare unseen data
unseen_df = pd.read_csv("unseen_df.csv")

# Encode categorical variables in the unseen data
label_encoder = LabelEncoder()
unseen_df["Gender"] = label_encoder.fit_transform(df["Gender"]) #
Male: 1, Female: 0
unseen_df["Marital_Status"] =
label_encoder.fit_transform(df["Marital_Status"])

```

```

# Define features (X_unseen) and target (y_unseen) for unseen data
X_unseen = unseen_df.drop(columns=["Purchase"])
y_unseen = unseen_df["Purchase"]

# Make predictions on unseen data
y_unseen_pred = classifier.predict(X_unseen)

# Evaluate the model's performance on unseen data
unseen_accuracy = accuracy_score(y_unseen, y_unseen_pred)
unseen_classification_report = classification_report(y_unseen,
y_unseen_pred)

unseen_accuracy, unseen_classification_report
(0.9992,
'
precision    recall  f1-score   support\n
0          1.00      1.00      1.00      4665\n
0.99      0.99      0.99      0.99      335\n
1.00      1.00      1.00      1.00      5000\n
weighted avg          1.00      1.00      1.00      5000\n
accuracy
macro avg          1.00      1.00      1.00      5000\n
nweighted avg          1.00      1.00      1.00      5000\n')

```

The decision tree classifier performed exceptionally well on the dataset, achieving an accuracy of 100% on the test set